1. Introduction and Research Context

The core problem being set up here is generalization failure - models trained on one dataset performing badly when deployed elsewhere. The healthcare framing makes this higher stakes than typical ML settings; wrong predictions can mean wrong treatment.

The review covers 15 sources, grouped into 5 methodological clusters. The clustering isn't just thematic - the argument is that each cluster reflects a different *theory* of what generalization failure actually is. That's a useful framing. Not just "here are papers on the same topic" but "here are fundamentally different ways of diagnosing the problem."

6 of the 15 get deeper critical treatment using CRAAP (Currency, Relevance, Authority, Accuracy, Purpose).

Main finding flagged upfront: the field is good at diagnosing the problem but hasn't produced a validated practical framework for predicting generalizability *before* deployment, or monitoring/adjusting after. Each cluster contributes something but none solves it fully. That's the stated gap.

Things to follow up on when reading the rest: how the 5 clusters are defined, whether the CRAAP integration actually works or feels forced, and whether the "gap" claim holds up once the individual papers are examined.